

# A Monte Carlo Approach for Power Estimation

Richard Burch, Farid N. Najm, *Member, IEEE*, Ping Yang, *Fellow, IEEE*, and Timothy N. Trick, *Fellow, IEEE*

**Abstract**—Excessive power dissipation in integrated circuits causes overheating and can lead to soft errors and/or permanent damage. The severity of the problem increases in proportion to the level of integration, so that power estimation tools are badly needed for present-day technology. Traditional simulation-based approaches simulate the circuit using test/functional input pattern sets. This is expensive and does not guarantee a meaningful power value. Other recent approaches have used probabilistic techniques in order to cover a large set of input patterns. However, they trade off accuracy for speed in ways that are not always acceptable. In this paper, we investigate an alternative technique that combines the accuracy of simulation-based techniques with the speed of the probabilistic techniques. The resulting method is *statistical* in nature; it consists of applying randomly-generated input patterns to the circuit and monitoring, with a simulator, the resulting power value. This is continued until a value of power is obtained with a desired *accuracy*, at a specified *confidence* level. We present the algorithm and experimental results, and discuss the superiority of this new approach.

## I. INTRODUCTION

EXCESSIVE power dissipation in integrated circuits causes overheating and can lead to soft errors and permanent damage. The severity of the problem increases in proportion to the level of integration. The advent of VLSI has led to much recent work on the estimation of power dissipation during the design phase, so that designs can be modified before manufacturing.

Perhaps the most significant obstacle in trying to estimate power dissipation is that the power is *pattern dependent*. In other words, it strongly depends on the input patterns being applied to the circuit. Thus the question "what is the power dissipation of this circuit?" is only meaningful when accompanied with some information on the circuit inputs.

A direct and simple approach of estimating power is to simulate the circuit. Indeed, several *circuit simulation* based techniques have appeared in the literature [1], [2]. Given the speed of circuit simulation, these techniques can not afford to simulate large circuits for long-enough input vector sequences to get meaningful power estimates. In order to simplify the problem and improve the speed, the power supply voltage is often assumed to be the same throughout the chip. Thus the power estimation problem is reduced to that of estimating the power supply *currents* that are drawn by the different circuit components. Fast timing or logic simulation can then be used to estimate these currents [3].

We call these approaches *strongly* pattern dependent because they require the user to specify *complete* information about the input patterns. Recently, other approaches have been proposed [4], [5] that only require the user to specify *typical* behavior at the circuit inputs using *probabilities*. These may be called *weakly* pattern dependent. With little computational effort, these techniques allow the user to cover a huge set of possible input patterns. However, in order to achieve good accuracy, one must model the correlations between internal node values, which can be very expensive. As a result, these techniques usually trade off accuracy for speed. The resulting loss of accuracy is a significant issue that may not always be acceptable to the user.

In this paper, we investigate an alternative approach that combines the accuracy of simulation-based approaches with the weak pattern dependence of probabilistic approaches. The resulting approach is *statistical* in nature; it consists of applying randomly generated input patterns to the circuit and monitoring, with a simulator, the resulting power value. This is continued until a value of power is obtained with a desired *accuracy*, at a specified *confidence* level. Since it uses a *finite* number of patterns to estimate the power, which really depends on the *infinite* set of possible input patterns, this method belongs to the general class of so-called *Monte Carlo* methods. A most attractive property of Monte Carlo techniques is that they are *dimension independent*, meaning that the *number* of samples required to make a good estimate is independent of the problem size. We will show that this property indeed holds for our approach (see Table IV in Section V).

Both [4] and [5] use probabilities to compute the power consumed by *individual gates*, which are then summed up to give the total power. In this context, it was observed in [5] that it would be too expensive to estimate the individual gate powers using a simulation with randomly generated inputs. The key to the efficiency of our new approach is that, if one monitors the *total* power directly during the random simulation, sufficient accuracy is obtained in much less time than is required to compute the individual gate powers. The excellent speed performance and the simplicity of the implementation make this a very attractive approach for power estimation.

An approach similar to this was independently proposed in [6], but the treatment there is not very rigorous and overlooks some important issues. Furthermore, no comparisons were performed with other approaches to show the superiority of the approach. In this paper, we present a rigorous treatment that provides the theoretical justification of this method. We also present experimental results of our implementation and compare it to probabilistic approaches.

Manuscript received May 5, 1992; revised August 27, 1992 and November 9, 1992.

R. Burch and P. Yang are with the Semiconductor Process & Design Center, Texas Instruments Inc., Dallas, TX 75265.

F. N. Najm and T. N. Trick are with the Electrical Engineering Department, University of Illinois at Urbana-Champaign Urbana, IL 61801.

IEEE Log Number 9206235.

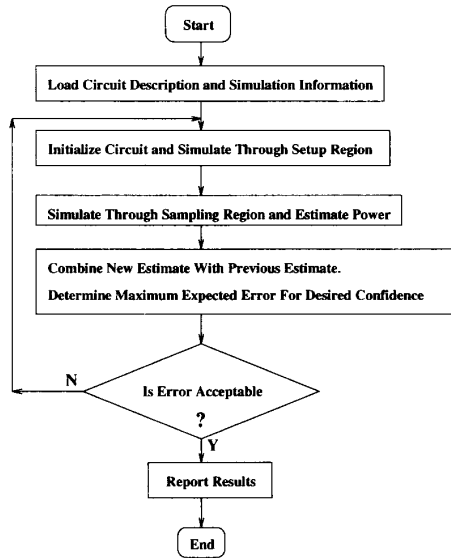


Fig. 1. Block diagram overview.

## II. OVERVIEW

In this section, we provide an overall view of our technique, and discuss its superiority to the probabilistic approaches previously proposed [4], [5].

### A. Overview of Monte Carlo Power Estimation

The block diagram in Fig. 1 gives an overall view of the technique. The *setup* and *sample* blocks are parts of the same logic simulation run, in which the input patterns are *randomly-generated*. The power value at the end of a sampling phase is noted and used to decide whether to *stop* the process or to do another setup-sample run. The decision is made based on the mean and standard deviation of the power values observed at the end of a number of successive iterations.

The power is found as the average value of the instantaneous power drawn throughout the sample phase, and *not* during the setup phase. However, the setup phase is a critical component of our approach, and serves two purposes.

- 1) In the beginning of the simulation run, the circuit does not switch as often as it typically would at a later time, when switching activity has had time to spread to all the gates. Thus the circuit is simulated until all nodes are switching at a stable rate during setup. This argument will be made more precise in Section III, where we also derive an exact value for the setup time.
- 2) The values of power observed at the end of successive sample intervals should be samples of independent random variables. This is required in order for the stopping criterion to be correct, and is guaranteed by restarting the random input waveforms at the beginning of the setup phase. The details are given in Section III.

Thus the setup phase guarantees that we are indeed measuring *typical* power, and ensures the *correctness* of the statistical stopping criterion.

### B. Comparison with Probabilistic Techniques

There are two distinct advantages of the Monte Carlo approach that make it an excellent choice for power estimation over probabilistic techniques. These are: 1) it achieves *desired* accuracy in reasonable time, avoiding the speed/accuracy tradeoff of probabilistic techniques and 2) the simplicity of the algorithm makes it very easy to implement in existing logic or timing simulation environments.

Probabilistic methods [4], [5] suffer from a speed/accuracy tradeoff because they must resolve the correlations between internal circuit nodes. If these correlations *are* taken into account, these methods can be very accurate. This, however, is computationally very expensive and impractical. As a result, fast implementations of these techniques are necessarily inaccurate. It is the aim of this paper to show that the proposed Monte Carlo method is very fast (solving circuits with thousands of gates in a matter of seconds) and also highly accurate (easily within 5% of the total power). Tables III and V compare the accuracy of Monte Carlo and probabilistic methods for power estimation.

We also should make the point that the accuracy level in our approach is predictable up-front: the program will work to achieve any level of accuracy desired by the user. Naturally, as higher accuracy is desired, the computational cost starts to increase. However, we will show in Section V that accuracy levels of 5% are easily and efficiently attainable.

## III. DETAILED APPROACH

This section describes the details of the approach. We start out with a rigorous formulation of the problem and show how it reduces to the well-known problem of *mean estimation* in statistics. We then discuss the stopping criterion, and the normality assumption required for it to work. We conclude with a discussion of the setup and sample phases and the applicability to sequential circuits.

### A. Problem Formulation

Consider a digital circuit with  $m$  internal nodes (gate outputs). Let  $x_i(t)$ ,  $t \in (-\infty, +\infty)$ , be the logic signal at node  $i$  and  $n_{x_i}(T)$  be the number of transitions of  $x_i$  in the time interval  $(-\frac{T}{2}, +\frac{T}{2}]$ . If, in accordance with [7], we consider only the contribution of the charging/discharging current components, the average power dissipated at node  $i$  during that interval is  $\frac{1}{2} V_{dd}^2 C_i \frac{n_{x_i}(T)}{T}$ , where  $C_i$  is the total capacitance at  $i$ . The total average power dissipated in the circuit during the same interval is

$$P_T = \frac{V_{dd}^2}{2} \sum_{i=1}^m C_i \frac{n_{x_i}(T)}{T}. \quad (1)$$

The power rating of a circuit usually refers to its average power dissipation over extended periods of time. We therefore define the *average power dissipation*  $P$  of the circuit as

$$P = \lim_{T \rightarrow \infty} P_T = \frac{V_{dd}^2}{2} \sum_{i=1}^m C_i \lim_{T \rightarrow \infty} \frac{n_{x_i}(T)}{T}. \quad (2)$$

The essence of our approach is to estimate  $P$ , corresponding to infinite  $T$ , as the mean of several  $P_T$  values, each measured over a finite time interval of length  $T$ . In order to see how this mean estimation problem comes about, we must consider a random representation of logic signals as follows.

Corresponding to every logic signal  $x_i(t)$ ,  $t \in (-\infty, +\infty)$ , we construct a stochastic process  $\mathbf{x}_i(t)$  as a family of the logic signals  $x_i(t + \tau)$ , where  $\tau$  is a random variable. This process has been called the companion process of  $x_i(t)$  in [5], [8], where the reader may find more details on its construction. For each  $\tau$ ,  $x_i(t + \tau)$  is a shifted copy of  $x_i(t)$ . Therefore, observing  $P_T$  for  $x_i(t + \tau)$  corresponds to measuring the power of  $x_i(t)$  over an interval of length  $T$  centered at  $\tau$ , rather than at 0. We can then talk of the random power of  $\mathbf{x}_i(t)$  over the interval  $(-\frac{T}{2}, +\frac{T}{2}]$ , to be denoted by

$$P_T = \frac{V_{dd}^2}{2} \sum_{i=1}^m C_i \frac{\mathbf{n}_{x_i}(T)}{T} \quad (3)$$

where  $\mathbf{n}_{x_i}(T)$  is now a random variable. It was shown in [5], [8] that  $\mathbf{x}_i(t)$  is stationary [12] so that, for any  $T$ , the expected average number of transitions per second is a constant:

$$D(x_i) = E\left[\frac{\mathbf{n}_{x_i}(T)}{T}\right] = \lim_{T \rightarrow \infty} \frac{\mathbf{n}_{x_i}(T)}{T} \quad (4)$$

where  $E[\cdot]$  denotes the expected value (mean) operator. In [5] and [8],  $D(x_i)$  was called the transition density of  $x_i(t)$ ; it is the average number of transitions per second, equal to twice the average frequency. As a result of (4),  $E[P_T]$  is the same for any  $T$ , and the average power can be expressed as a mean:

$$P = E[P_T]. \quad (5)$$

Thus the power estimation problem has been reduced to that of mean estimation, which is a frequently encountered problem in statistics.

In order to apply the above theory, we must ensure that the signals  $x_i(t)$  observed throughout the  $(-\frac{T}{2}, +\frac{T}{2}]$  interval are samples of the stationary processes  $\mathbf{x}_i(t)$ . This requirement will be addressed in Section III-D.

### B. Stopping Criterion

To determine stopping criteria, some assumptions about the distribution of  $P_T$  must be made. Therefore, let us assume that  $P_T$  is normally distributed for any  $T$ . The theoretical justification and experimental evidence for this assumption will be discussed in Section III-C, and Section IV will discuss the applicability of Monte Carlo methods when this assumption is invalid. Suppose also that we perform  $N$  different simulations of the circuit, each of length  $T$ , and form the sample average  $\eta_T$  and sample standard deviation  $s_T$  of the  $N$  different  $P_T$  values found. Therefore, we have  $(1 - \alpha) \times 100\%$  confidence that  $|\eta_T - E[P_T]| < t_{\alpha/2} s_T / \sqrt{N}$ , where  $t_{\alpha/2}$  is obtained from the  $t$  distribution [9] with  $(N - 1)$  degrees of freedom. This result can be rewritten as

$$\frac{|P - \eta_T|}{\eta_T} < \frac{t_{\alpha/2} s_T}{\eta_T \sqrt{N}}. \quad (6)$$

Therefore, for a desired percentage error  $\epsilon$  in the power estimate, and for a given confidence level  $(1 - \alpha)$ , we must simulate the circuit until

$$\frac{t_{\alpha/2} s_T}{\eta_T \sqrt{N}} < \epsilon. \quad (7)$$

We can use this relation to illustrate the important dimension independence property of this approach, common to most Monte Carlo methods, as follows. If  $N^*$  is the (smallest) number of iterations that satisfies (7), then,

$$N^* \approx \left( \frac{t_{\alpha/2} s_T}{\epsilon \eta_T} \right)^2. \quad (8)$$

By dimension independence, we mean that  $N^*$  should be roughly independent of the circuit size (number of nodes). In (7),  $t_{\alpha/2}$  is a small number, typically between 2.0 and 5.0, and  $\epsilon$  is a constant. We, therefore, look to the ratio  $s_T^2 / \eta_T^2$  to learn of the general behavior of  $N^*$ . There is very little one can say in general about this ratio. Nevertheless, and in view of (3), it is instructive to consider the following. If  $\mathbf{y} = \sum_{i=1}^m \mathbf{x}_i$  is the sum of  $m$  independent identically distributed (IID) random variables, then  $\sigma_y^2 / \mu_y^2 = (\sigma_{x_i}^2 / \mu_{x_i}^2) \times (1/m)$ , where  $\mu$  and  $\sigma^2$  denote the mean and variance. Thus if the  $C_i \mathbf{n}_{x_i}(T) / T$  terms of (3) are iid, then  $N^*$  should decrease with circuit size. Even when the  $\mathbf{x}_i$ 's are not independent, we have  $\sigma_y^2 / \mu_y^2 \leq (\sigma_{x_i}^2 / \mu_{x_i}^2)$ , a constant, which suggests that  $N^*$  should typically decrease or remain constant with increasing circuit size. This is indeed the observed behavior in Table IV.

An important consequence of this result is that, since each iteration of the Monte Carlo approach takes roughly linear time (in the size of the circuit), then the overall process should also take linear time. Probabilistic methods that do not take correlation into account also depend linearly on circuit size. However, if correlation is taken into account in order to improve the accuracy, their dependence is frequently super-linear.

In order to use the stopping criterion in practice, we must ensure that the observed  $P_T$  values are samples from independent  $P_T$  random variables. This requirement will be addressed in Section III-D.

### C. Normality

It cannot be proven that  $P_T$  is normally distributed for finite  $T$ ; however,  $P_T$  can frequently be approximated as normal. This section will discuss sufficient conditions for the normality of  $P_T$  that make the approximation sensible, and show experimental evidence that the approximation is good on a variety of combinational benchmark circuits.

A sufficient condition for the normality of  $P_T$  is that (i)  $m$  is large and (ii)  $\mathbf{n}_{x_i}(T) / T$  are independent. This is true under fairly general conditions irrespective of the individual  $\mathbf{n}_{x_i}(T) / T$  distributions (see [10, pp. 188–189]), and for any value of  $T$ .

Another sufficient condition that holds even for small  $m$ , but for large  $T$ , is as follows. If (i) the consecutive times between identical transitions of  $\mathbf{x}_i(t)$  are independent (which, using renewal theory (see [11, pp. 62–63]), means that  $\mathbf{n}_{x_i}(T) / T$




Fig. 2. Normal scores plot for the ISCAS-85 circuits.

is normally distributed for large  $T$  and (ii) the  $\mathbf{n}_{x_i}(T)/T$  are independent (so that they are also jointly normal (see [12], p. 126) for large  $T$ ) then  $P_T$  is normal for large  $T$  (see [12, p. 144]).

To the extent that these conditions are approximately met in practice, the power should be approximately normal. In both cases, the condition that the  $\mathbf{n}_{x_i}(T)/T$  are independent is not met; however, if a significant subset of the  $\mathbf{n}_{x_i}(T)/T$  are approximately independent, then it is reasonable to expect that the power should be approximately normal. We have found that for a number of benchmark digital circuits [13], the normality assumption is very good, as shown in the *normal scores plots* [9] in Fig. 2. The plot for each circuit corresponds to 1000 evaluations of the average power over a 2.5- $\mu$ s interval. Each evaluation covered an average of 50 transitions per primary input. It cannot be proven that the power of all circuits is approximately normal. The consequences of deviations from normality are discussed in Section IV.

#### D. Setup and Sample

This section deals with the mechanics of how the input patterns are to be generated, when to start and stop measuring a  $P_T$  value, and how different  $P_T$  values should be obtained. We start by observing that, by stationarity of  $\mathbf{x}_i(t)$ , the (finite) intervals of width  $T$ , over which the  $P_T$  values will be measured, need not be centered at the origin. A  $P_T$  value may be obtained from any interval of width  $T$ , henceforth called a *sampling interval*. However, the following two requirements must be met:

- 1) (i) Throughout a sampling interval, the signals  $x_i(t)$  must be samples of the *stationary* processes  $\mathbf{x}_i(t)$ .
- 2) (ii) The different  $P_T$  samples must be samples from *independent*  $P_T$  random variables.

We will now describe a simulation process that guarantees both of these requirements. Suppose that the circuit primary inputs are at 0 from  $-\infty$  to time 0, and then become samples of the stationary processes  $\mathbf{x}_i(t)$  in positive time. Consider a primary input driving an inverter with delay  $t_d$ . Since its input is a stationary process for  $t \geq 0$ , its output must be stationary for  $t \geq t_d$ . By using a simplified timing model for every gate




Fig. 3. Successive setup and sample phases.

as in [5], we can repeat this argument enough times to obtain the following conclusion: If the maximum delay (along any path) from the primary inputs to node  $i$  is  $T_{max,i}$ , then the process  $\mathbf{x}_i(t)$  becomes stationary for  $t \geq T_{max,i}$ .

If the maximum delay (along any path) in the circuit is  $T_{max} = \max_i(T_{max,i})$ , then the sampling interval may start only after  $t \geq T_{max}$ . This guarantees that requirement (i) is met. From that time onwards, all internal processes are stationary, and the circuit is in (probabilistic) steady state. We will call the time interval from 0 to  $T_{max}$  the *setup* phase. Intuitively, the circuit needs to be simulated until all nodes are switching at stable rates before a reliable sample of power may be taken and, as we have shown, the minimum time required to achieve that is  $T_{max}$ . Finding  $T_{max}$  in combinational circuits is straightforward; the case of sequential (feedback) circuits is discussed in Section III-E.

In order to guarantee requirement (ii), we simply *restart* the simulation (with an empty event queue) at the beginning of every setup phase. As a result, the time axis is divided into successive regions of setup and sampling, as shown in Fig. 3.

We have discussed the length of the setup region. Now, consider the length of the sampling region,  $T$ . Two factors must be considered: the normality approximation and the overall simulation time. The length of the sampling region can affect the error in the normality approximation. The minimum value of  $T$  that yields approximately normal power is heavily dependent upon the circuit and can range from nearly zero to nearly infinity. To minimize the simulation time, minimize  $N \times (T_{max} + T)$ .  $N$  depends upon  $s_T$ , and  $s_T$  decreases as  $T$  increases. The speed of this decrease is heavily dependent upon the circuit. These two factors indicate that no general solution for optimum  $T$  exists. We chose a value for  $T$  (see results) that worked well experimentally.

The only remaining task is to describe how the inputs are to be generated. This has to be done in such a way that the input processes, after the start of every setup phase, are independent of the past. This can be done as follows, for every input signal  $x_i$ . At the beginning of a setup phase, we use a random number generator to select a logic value for  $x_i$ , with appropriate signal probability  $P(x_i)$ . We then use another random number generator to decide how long  $x_i$  stays in that state before switching. This must assume some distribution for the duration of stay in that state. Once  $x_i$  has switched, we use another random number generator to decide how long it will stay in the other state, again using some distribution. Let  $F_{x_i}^1(t)$  be the distribution of times spent in the 1 state, and  $F_{x_i}^0(t)$  be that of the 0 state. Since computer implementations of random number generators produce sequences of independent random variables, independence between the successive sampling phases is guaranteed.

The starting signal probability  $P(x_i)$  and distributions  $F_{x_i}^1(t)$  and  $F_{x_i}^0(t)$  should be supplied by the user. In fact, these parameters represent the way in which the approach is weakly pattern dependent. They also provide the mechanism by which the user can specify any information about typical behavior at the circuit inputs. In order to simplify the user interface, our current implementation does not require the user to actually specify distributions. Rather, we require only two parameters: the average time that an input is high, denoted by  $\mu_{x_i}^1$ , and the average time that it is low, denoted by  $\mu_{x_i}^0$ . Based on this, it can be shown [5, 8] that  $P(x_i) = \mu_{x_i}^1 / (\mu_{x_i}^1 + \mu_{x_i}^0)$  and  $D(x_i) = 2 / (\mu_{x_i}^1 + \mu_{x_i}^0)$ . As for the distributions, our implementation uses *exponential* distributions [12] so as to allow the comparisons with probabilistic methods to be given in Section V. We emphasize, however, that the stopping criterion and the overall Monte Carlo algorithm are valid for *any* distribution. In fact, in our implementation, the choice of distribution can be easily modified by the user. Furthermore, experiments with different input distributions have shown that the average power of many circuits is not sensitive to the form of the distributions.

Previously, we have assumed that the primary inputs are independent. Because of the expense of handling correlation, this is necessary for probabilistic methods. If the correlation can be expressed as conditional probabilities, then logical waveforms can be generated that take the correlation into account and Monte Carlo methods can be used to estimate the power effectively.

#### E. Sequential Circuits

The Monte Carlo method presented in this paper is valid for both combinational and sequential circuits. The only aspect of the problem that is specific to sequential circuits is the computation of the setup time  $T_{max}$ . Strictly speaking, since  $T_{max}$  is the longest delay along *any* path, then  $T_{max} = \infty$  for sequential circuits. Recall that it is *sufficient* to wait for  $T_{max}$  before starting a sampling interval in order to guarantee stationarity. It is not clear, however, whether that condition is also necessary. Heuristics must be developed to estimate the setup lengths for sequential circuits. The quality of a heuristic can be tested by examining the expected power for sampling regions of different lengths. If the expected power is constant, then the heuristic does a good job of predicting the length of the setup region. Table I shows that this is true for combinational circuits, which agrees with our assertion in Section III that the power would be stationary after  $T_{max}$ . Future implementations of this approach will include heuristics to allow it to handle sequential circuits.

#### IV. DEVIATIONS FROM NORMALITY

Monte Carlo method can be applied to circuits that have nonnormal power distributions without adversely affecting the accuracy of the results. In cases of severe deviations from normality, some modifications of the basic approach may be required. It is important to note at the outset that the normality assumption was required only to formulate the stopping criterion. Given enough samples, one ultimately

TABLE I  
IF THE SETUP REGION IS CHOSEN CORRECTLY, THEN THE PROCESS IS STATIONARY AND THE EXPECTED POWER IS THE SAME FOR ANY SAMPLING REGION. THIS IS ILLUSTRATED FOR COMBINATIONAL CIRCUITS WITH SAMPLING REGIONS OF 625 ns, 1.25  $\mu$ s AND 2.5  $\mu$ s

Circuit Name	Power		
	T = 625 ns	T = 1.25 $\mu$ s	T = 2.5 $\mu$ s
c432	1.13 mW	1.13 mW	1.12 mW
c499	2.04 mW	2.04 mW	2.05 mW
c880	2.75 mW	2.74 mW	2.75 mW
c1355	5.45 mW	5.45 mW	5.45 mW
c1908	9.23 mW	9.25 mW	9.22 mW
c2670	10.78 mW	10.79 mW	10.80 mW
c3540	14.57 mW	14.62 mW	14.64 mW
c5315	23.14 mW	23.10 mW	23.10 mW
c6288	70.38 mW	70.36 mW	70.32 mW
c7552	37.57 mW	37.50 mW	37.50 mW

converges to the desired power value, *whatever* the power distribution. This is true because of (5) and the *strong law of large numbers* (see [11, p. 26]). Furthermore, we are only concerned with deviations from normality for small values of  $T$ . For large  $T$ , and since  $P_T$  tends to a constant as  $T \rightarrow \infty$ , the variance of  $P_T$  goes to 0, and its distribution must become bell-shaped, approaching a normal distribution.

Let  $P_T$  have any distribution. The *central limit theorem* [9] implies that if  $N$  samples of  $P_T$  are averaged, then the distribution of the averages approaches normal as  $N \rightarrow \infty$ . Convergence criteria exist that are similar to those derived for normally distributed power. For a desired percentage error  $\epsilon$  in the power estimate and for a given confidence level  $(1 - \alpha)$ , we must simulate the circuit until

$$\frac{z_{\alpha/2} \sigma_T}{\eta_T \sqrt{N}} < \epsilon$$

$\sigma_T$ , the standard deviation of  $P_T$ , is not known. When  $\sigma_T$  is replaced with the sample standard deviation  $s_T$ , error is introduced and some assumption about the distribution of  $P_T$  is necessary. We will assume that when  $\frac{t_{\alpha/2} s_T}{\eta_T \sqrt{N}} < \epsilon$ , then  $\frac{z_{\alpha/2} \sigma_T}{\eta_T \sqrt{N}} < \epsilon$ . This is equivalent to assuming that  $P_T$  is normal; however, it allows us to consider the consequences of nonnormal distributions. If  $s_T$  is large, then many samples must be taken before convergence, and with a large enough number of samples,  $s_T \approx \sigma_T$ . If  $s_T$  is small enough to allow convergence for small  $N$ , then it is likely that  $\sigma_T$  is also small. If the distribution has a large concentration of values around one power and a small number of values at a much higher (or lower) power, then errors may result because  $s_T$  is likely to be much smaller than  $\sigma_T$ . The rest of this section considers various nonnormal distributions, including a *bimodal* distribution that illustrates the error in the normality assumption. The final part of this section suggests an approach to identify the types of circuits that would cause an error.

Circuits do exist that have nonnormal power for small  $T$ . An example would be a circuit with an enable signal whose value strongly affects the power drawn by the whole circuit. When the enable signal is low, the circuit would have one power distribution, and when it is high it would have a different




Fig. 4. A bimodal distribution.




Fig. 5. A circuit with enable.

distribution. If these two distributions had different means, then over a small  $T$  interval the overall distribution would not be normal, but would be a so-called *bimodal* distribution, as shown in Fig. 4. When each of the two distributions is normal, we will refer to the overall distribution as a *double normal*.

As a concrete example, consider the simple XOR circuit in Fig. 5. The enable signal allows the output stage to switch, drawing much more power than otherwise. If the transition density (see equation (4)) at the enable line is low compared to the other two inputs, then, over a short time interval, only one of the two modes of operation would be observed. Using transition densities of  $2e7$  on the inputs and  $2e5$  on the enable line, a sampling interval of  $T = 2.5\mu\text{sec}$ , and with 11 000 samples, we get the histogram shown in Fig. 6.

This is one of many ways in which the distribution can deviate from normality. We have also considered two other ways in which the distribution can be distorted. We will refer to distributions with elongated tails, as shown in Fig. 7(a), as *tailed normal* distributions. Those with chopped tops, as in Fig. 7(b), will be called *chopped normal* distributions.

We have examined the performance of our stopping criteria for each of the above varieties of distorted normals. The nonnormal power values were artificially obtained from customized random number generators. The parameters for the stopping criterion were set to 5% accuracy ( $\epsilon = 0.05$ ) with 99% confidence ( $\alpha = 0.01$ ).

Since the distributions were not normal, one would expect the resulting accuracy to be somewhat worse than 5%. In all but a few cases, better than 10% accuracy was achieved with 99% confidence. The only examples that showed worse




Fig. 6. Power distribution for the XOR circuit.




Fig. 7. Distortions of the distributions. (a) Tailed normal. (b) Chopped normal.

than 10% accuracy were distributions with very long tails, and double normal distributions with widely separated means. Even the distributions with very long tails had better than 15% accuracy with 99% confidence. The double normal distributions, however, had very large errors if the first few samples were all centered around one hump of the distribution. When this occurred, the stopping criteria erroneously terminated the simulation. This is exactly as anticipated.

In the cases of the drastic double normal distributions, we feel that they can be treated as follows. When a node has a high fanout, making it a potential cause of the double normal, then the length of each sampling region should be changed so that that node transitions several times in a sampling region. This will prevent the problems associated with an enable signal and should prevent problems with any double normal distributions.

Having said all this, and before leaving this section, it is important to reiterate that the normality assumption holds very well for all the benchmark circuits that we have considered, as discussed in Section III and shown in Fig. 2.

## V. EXPERIMENTAL RESULTS

The Monte Carlo methods presented in this paper were implemented based on a simple variable delay logic simulator. This program will be referred to as McPOWER. The test circuits to be used in this section are the benchmarks presented at ISCAS in 1985 [13]. These circuits are combinational logic

TABLE II  
THE ISCAS-85 BENCHMARK CIRCUITS

Circuit	#inputs	#outputs	#gates
c432	36	7	160
c499	41	32	202
c880	60	26	383
c1355	41	32	546
c1908	33	25	880
c2670	233	140	1193
c3540	50	22	1669
c5315	178	123	2307
c6288	32	32	2406
c7552	207	108	3512

circuits and Table II presents the number of inputs, outputs, and gates in each.

We will compare the performance of McPOWER to that of probabilistic methods and substantiate the claims of Section II-B that McPOWER has better accuracy and competitive simulation times. DENSIM [5], [8] is an efficient probabilistic simulation program that gives the *average switching frequency* (called *transition density* in [5], [8]) at every circuit node. These density values can be used to give an estimate of the total power dissipation. DENSIM does not take into account the correlation between internal circuit nodes. While it is known that this causes inaccuracy in the density values, it is prohibitively expensive to take all correlation into account for large circuits.

Table III compares the performance of DENSIM, when used to estimate total power, to that of McPOWER. In both programs, every primary input had a signal probability of 0.5 and a transition density of  $2e7$  transitions per second (corresponding to an average frequency of 10 MHz). For McPOWER, a maximum error of 5% with 99% confidence was specified. As mentioned in Section III, McPOWER performs one long simulation that is broken into setup and sampling regions. The delays of the circuit determine the length of each setup region; however, the length of a sampling region is specified by the user. For Table III, the sampling region was set to  $2.5 \mu s$ , which allows an average of 50 transitions per sampling interval on each input. The column labeled LOGSIM gives our best estimates of the power dissipation of these circuits, obtained from one very long logic simulation for each circuit. As seen from the table, McPOWER is consistently and highly accurate, while DENSIM has significant errors for some circuits. Although DENSIM is frequently faster, McPOWER's reliable accuracy makes it a more attractive approach for power estimation.

The typical, rapid convergence of McPOWER is illustrated in Fig. 8. The figure shows the power from three different iterations converging to the average power for c6288, one of the most complex ISCAS circuits. A similar plot is shown for c5315 in Fig. 9.

Care must be taken in drawing conclusions from a single run of McPOWER. Since it uses random input vectors, the speed of convergence and the error in the power estimate depend on the initialization of the random number generator. This is

TABLE III  
POWER AND TIME RESULTS FOR THE ISCAS-85 CIRCUITS. TIME IS IN CPU SECONDS ON A SUN SPARCSTATION1. McPOWER IS BASED ON 5% ERROR, 99% CONFIDENCE AND,  $2.5 \mu s$  SAMPLING REGION

Circuit Name	Power		CPU Time	
	DENSIM	LOGSIM	DENSIM	McPower
c432	0.974 mW	1.165 mW	1.17 mW	0.7 s (3.6X)
c499	1.977 mW	2.048 mW	2.13 mW	0.8 s (2.8X)
c880	2.086 mW	2.829 mW	2.81 mW	1.4 s (2.7X)
c1355	3.695 mW	5.735 mW	5.65 mW	1.9 s (1.9X)
c1908	5.154 mW	9.734 mW	9.77 mW	3.1 s (1.8X)
c2670	7.319 mW	11.438 mW	11.24 mW	4.5 s (1.6X)
c3540	9.235 mW	15.328 mW	15.25 mW	5.8 s (2.1X)
c5315	15.471 mW	24.102 mW	23.66 mW	8.5 s (2.6X)
c6288	31.941 mW	78.883 mW	75.53 mW	40.4 s (5.4X)
c7552	23.156 mW	40.006 mW	38.78 mW	24.7 s (2.0X)




Fig. 8. McPOWER convergence results for c6288.

illustrated in Table IV, which shows the statistics obtained from 1000 McPOWER runs. The minimum, maximum, and average number of iterations required per run for 5% accuracy with 99% confidence are given. Notice that the average number of iterations required to converge does *not* increase with the circuit size. This confirms the *dimension independence* property of this approach which, as pointed out in the introduction, is a common feature of Monte Carlo methods. Also shown in the table are the percentage number of runs for which the error was greater than 5%, which, as expected, is less than 1% in all cases.

Even smaller execution times are possible if the desired accuracy and confidence levels are relaxed. Table V compares DENSIM to a single run of McPOWER with 95% confidence, 20% accuracy, and a sampling region of 625 ns. As in

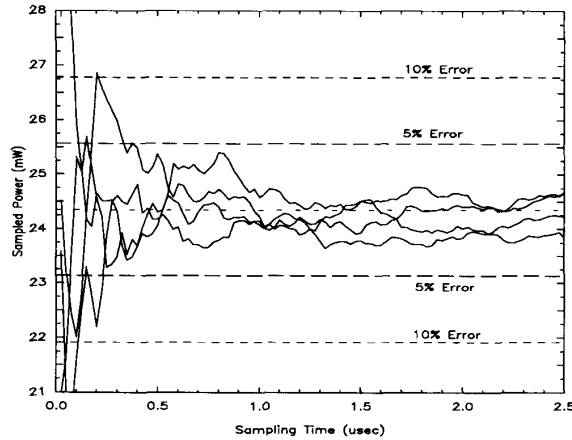


Fig. 9. McPOWER convergence results for c5315.

TABLE IV  
STATISTICS FROM 1000 McPOWER RUNS

Name	Min	Max	Avg	% > 5% Err	Avg CPU Time
c432	3	15	8.0	0.9%	3.2 s
c499	3	7	3.9	0.0%	2.9 s
c880	3	14	6.7	0.4%	3.9 s
c1355	3	7	4.0	0.0%	4.8 s
c1908	3	11	5.6	0.3%	10.3 s
c2670	3	10	5.2	0.1%	12.4 s
c3540	3	13	6.0	0.0%	18.3 s
c5315	3	7	4.1	0.0%	22.4 s
c6288	3	6	3.8	0.0%	50.4 s
c7552	3	10	5.5	0.0%	44.4 s

Table III, each input has a signal probability of 0.5 and a transition density of  $2e7$  transitions per second. With these parameters, McPOWER shows competitive speed and still exhibits superior accuracy.

## VI. CONCLUSIONS

In this paper, we have considered the problem of power estimation of a VLSI circuit. We assumed that standby power was negligible; although, standby power can be treated in a very similar manner. Furthermore, we only considered the charging/discharging current since this is the most significant component [7]. Other sources of transient current can occur and are handled identically to the charging/discharging current. They were not considered to avoid complication without additional insight.

We have presented a Monte Carlo based power estimation method. Randomly generated input waveforms are applied to the circuit using a logic/timing simulator and the cumulative value of total power is monitored. The simulation is stopped when sufficient accuracy is obtained with specified confidence. The statistical stopping criterion was discussed, along with experimental results from our prototype implementation McPOWER.

TABLE V  
POWER AND TIME RESULTS FOR THE ISCAS-85 CIRCUITS. TIME IS IN CPU SECONDS ON A SUN SPARCSTATION1. McPOWER IS BASED ON 20% ERROR, 95% CONFIDENCE AND, 625 NS SAMPLING REGION

Circuit Name	Power		CPU Time	
	DENSIM	LOGSIM	DENSIM	McPower
c432	0.974 mW	1.165 mW	1.10 mW	0.7 s (1.0X)
c499	1.977 mW	2.048 mW	2.04 mW	0.8 s (1.1X)
c880	2.086 mW	2.829 mW	2.91 mW	1.4 s (0.7X)
c1355	3.695 mW	5.735 mW	5.41 mW	1.9 s (0.8X)
c1908	5.154 mW	9.734 mW	9.27 mW	3.1 s (0.7X)
c2670	7.319 mW	11.438 mW	10.83 mW	4.5 s (0.6X)
c3540	9.235 mW	15.328 mW	14.28 mW	5.8 s (0.8X)
c5315	15.471 mW	24.102 mW	22.65 mW	8.5 s (0.8X)
c6288	31.941 mW	78.883 mW	71.48 mW	7.5 s (1.3X)
c7552	23.156 mW	40.006 mW	37.88 mW	12.4 s (0.7X)

We have shown that Monte Carlo methods are, in general, better than probabilistic methods for the estimation of power since they achieve superior accuracy with comparable speeds. They are also easier to implement and can be added to existing timing or logic simulation tools. Furthermore, the accuracy can be specified up-front with any desired confidence. Although the type of circuit may affect the amount of power drawn or the number of samples needed to converge, it will not affect the accuracy or reliability of these methods.

Feedback circuits present a severe problem for probabilistic methods. Monte Carlo methods are based on simple timing or logic simulation techniques and, therefore, experience very few difficulties with feedback circuits. The only unresolved problem is to determine the length of the setup region, but we feel that good heuristics can be developed for this. Future research will focus on developing such heuristics, thus generalizing the Monte Carlo technique to handle any logic circuit.

Although we have clearly demonstrated the superiority of Monte Carlo methods for power estimation, it is not clear that they will be better than probabilistic methods for other applications, such as estimating the power supply current waveforms. Future research will be aimed at exploring this and other applications of the Monte Carlo approach.

## REFERENCES

- [1] S. M. Kang, "Accurate simulation of power dissipation in VLSI circuits," *IEEE J. Solid-State Circuits*, vol. SC-21, pp. 889-891, Oct. 1986.
- [2] G. Y. Yacoub and W. H. Ku, "An accurate simulation technique for short-circuit power dissipation based on current component isolation," *IEEE Int. Symp. on Circuits and Systems*, pp. 1157-1161, 1989.
- [3] A.-C. Deng, Y.-C. Shiao, and K.-H. Loh, "Time domain current waveform simulation of CMOS circuits," *IEEE Int. Conf. on Computer-Aided Design*, Santa Clara, CA, pp. 208-211, Nov. 7-10, 1988.




- [4] M. A. Cirit, "Estimating dynamic power consumption of CMOS circuits," *IEEE Int. Conf. on Computer-Aided Design*, pp. 534-537, Nov. 9-12, 1987.
- [5] F. Najm, "Transition density, a stochastic measure of activity in digital circuits," *28th ACM/IEEE Design Automation Conf.*, San Francisco, CA, pp. 644-649, June 17-21, 1991.
- [6] C. M. Huizer, "Power dissipation analysis of CMOS VLSI circuits by means of switch-level simulation," *IEEE European Solid State Circuits Conf.*, pp. 61-64, Grenoble, France, 1990.
- [7] H. J. M. Veendrick, "Short-circuit dissipation of static CMOS circuitry and its impact on the design of buffer circuits," *IEEE J. Solid-State Circuits*, vol. SC-19, pp. 468-473, Aug. 1984.
- [8] F. Najm, "Transition density, a new measure of activity in digital circuits," *IEEE Trans. Computer-Aided Design*, vol. 12, pp. 310-323, Feb. 1993.
- [9] I. R. Miller, J. E. Freund, and R. Johnson, *Probability and Statistics for Engineers*. Englewood Cliffs, NJ: Prentice Hall, 1990.
- [10] A. Hald, *Statistical Theory with Engineering Applications*. New York: Wiley, 1952.
- [11] S. M. Ross, *Stochastic Processes*. New York: Wiley, 1983.
- [12] A. Papoulis, *Probability, Random Variables, and Stochastic Processes*, 2nd edition. New York: McGraw-Hill, 1984.
- [13] F. Brglez and H. Fujiwara, "A neutral netlist of 10 combinational benchmark circuits and a target translator in Fortran," *IEEE Int. Symp. on Circuits and Systems*, pp. 695-698, June 1985.



**Richard Burch** received the B.S. degree in physics and the M.S.E.E. degree from the University of Oklahoma, in 1984 and 1986, respectively. He is currently working toward the Ph.D. degree in electrical engineering at the University of Illinois at Urbana-Champaign, Urbana.

In 1984, he joined Texas Instruments' Semiconductor Process and Design Center as a Member of the Technical Staff. His current interests include research into circuit simulation, power and current estimation, and electromigration failure prediction.

In 1992 Mr. Burch received the award for the best paper published in the *IEEE TRANSACTIONS ON COMPUTER-AIDED DESIGN*.




**Farid N. Najm** (S'85-M'89) received the B.E. degree (with distinction) in electrical engineering from the American University of Beirut, Lebanon, in 1983, and the M.S. and Ph.D. degrees in electrical and computer engineering from the University of Illinois at Urbana-Champaign, in 1986 and 1989, respectively.

He was with the General Swedish Electric Company (ASEA) in Västerås, Sweden, in 1982, and was a teaching assistant at AUB in 1983. He later worked as an electronics engineer with AUB from 1983 to 1984, and held a training position with the University of Warwick,

England, in 1984. While at the University of Illinois, in 1985-1989, he was a research assistant with the Coordinated Science Laboratory, and held training positions with the VLSI Design Laboratory at Texas Instruments, Dallas, TX. In July 1989, he joined Texas Instruments as a Member of Technical Staff with the Semiconductor Process and Design Center. In August 1992, he became an Assistant Professor at the Department of Electrical and Computer Engineering, University of Illinois at Urbana-Champaign. His research interests include design for reliability, synthesis of reliable and low-power VLSI, probabilistic circuit analysis and simulation, timing analysis, test generation, and circuit and timing simulation.

**Ping Yang** (M'81-SM'86-F'91), for a photograph and biography please see page 6 of this issue.



**Timothy N. Trick** (S'63-M'65-SM'73-F'79) received the B.S. degree in electrical engineering from the University of Dayton, in 1961, and the M.S. and Ph.D. degrees from Purdue University, Lafayette, IN, in 1962 and 1966, respectively.

Since 1965 he has been with the University of Illinois at Urbana-Champaign, where he currently Professor and Head of the Department of Electrical and Computer Engineering. He has served as the Director of the Coordinated Science Laboratory at the University of Illinois from 1984 to 1986. In

the summers of 1970 and 1971 he was an ASEE-NASA Summer Faculty Fellow at the Marshall Space Flight Center, Huntsville, AL. During 1973 to 1974 he was Visiting Associate Professor with the Department of Electrical Engineering and Computer Sciences, University of California, Berkeley. He was a member of the IEEE Board of Directors 1986 to 1989, IEEE Vice-President for Publications, 1988 to 1989, President of the IEEE Circuits and Systems Society, 1979, and served on numerous technical committees and editorial boards. His research interests are computational methods for circuit analysis and design, integrated circuits, and analog and digital signal processing.

In 1976 Dr. Trick received the Guillemin-Cauer Award for best paper published in the *IEEE TRANSACTIONS ON CIRCUITS AND SYSTEMS*, the IEEE Centennial Award in 1984, and in 1987 he received the IEEE Circuits and Systems Society Meritorious Award. He is a member of AAAS.